

Sanjay Kariyappa

Last Updated on 14th June 2026

<https://sanjaykariyappa.github.io>
sanjaykariyappa@gmail.com | 678.650.5017 | Santa Clara, CA

SUMMARY

Sr. Research Scientist at NVIDIA. Interested in building secure and efficient AI.

EDUCATION

GEORGIA TECH

PHD IN ECE
Dec 2022 | Atlanta, GA
GPA: 4.0 / 4.0

GEORGIA TECH

MS IN ECE
Dec 2014 | Atlanta, GA
GPA: 4.0 / 4.0

SRI JAYACHAMARAJENDRA COLLEGE OF ENGINEERING

BS IN ECE
June 2013 | Mysore, India
GPA: 3.78 / 4.0

LINKS

Github:// [sanjaykariyappa](#)
LinkedIn:// [sanjay-kariyappa](#)
Twitter:// [@sanjayatwork](#)

WORK EXPERIENCE

NVIDIA | SR. RESEARCH SCIENTIST
Nov 2024 – present | Santa Clara, CA
Agent security and efficiency research.

JP MORGAN CHASE | SR. AI RESEARCH ASSOCIATE
Feb 2023 – Nov 2024 | Palo Alto, CA
Explainable AI research.

META | AI RESEARCH INTERN (FAIR)
May 2022 - Aug 2022 | Boston, MA
Developed scalable attack on federated learning. (paper)

FACEBOOK | SOFTWARE ENGINEERING INTERN
May 2021 - Aug 2021, May 2020 - Aug 2020 | Menlo Park, CA
Semi-supervised learning to improve conversion prediction models.

IBM | RESEARCH INTERN
May 2019 – Aug 2019 | San Jose, CA
Noise-Resilient DNNs for PCM-based analog AI hardware. (paper)

PUBLICATIONS

Where Instruction Hierarchy Breaks: Diagnosing and Repairing Failures in Reasoning Language Models [S. Kariyappa, G.E. Suh](#)

SideQuest: Model-Driven KV Cache Management for Long-Horizon Agentic Reasoning [S. Kariyappa, G.E. Suh](#)

Architecting Secure AI Agents: Perspectives on System-Level Defenses Against Indirect Prompt Injection Attacks C. Xiang, D. Zagieboylo, S. Ghosh, [S. Kariyappa](#), K Greshake, H. Xiao, C. Xiao, G.E. Suh

Mitigating Indirect Prompt Injection via Instruction-Following Intent Analysis M. Kang, C. Xiang, [S. Kariyappa](#), C. Xiao, G.E. Suh

Stronger Enforcement of Instruction Hierarchy via Augmented Intermediate Representations [S. Kariyappa, G.E. Suh](#)

Interpretable LLM-based Table Question Answering [TMLR 2025] G. Nguyen, I Brugere, S Sharma, [S. Kariyappa](#), A. Nguyen, F. Lécué

Progressive Inference: Explaining Decoder-Only Sequence Classification Models Using Intermediate Predictions [ICML 2024] [S. Kariyappa](#), F. Lécué, S. Mishra, C. Pond, D. Magazzeni, M. Veloso

SHAP@k: Efficient and PAC Identification of Top-k Features [AAAI 2024 - Oral] [S. Kariyappa](#), L. Tsepenekas, F. Lécué, D. Magazzeni

Cocktail Party Attack: Breaking Aggregation-Based Privacy in Federated Learning using Independent Component Analysis [ICML 2023] [S. Kariyappa](#), C. Guo, K. Maeng, W. Xiong, Ed Suh, M. K. Qureshi, H. S. Lee

ExPLOit: Extracting Private Labels in Split Learning [SaTML 2023] [S. Kariyappa](#), M. K. Qureshi

Bounding the Invertibility of Privacy-preserving Instance Encoding using Fisher Information [NeurIPS 2023] K. Maeng, C. Guo, [S. Kariyappa](#), Ed Suh

MAZE: Data-Free Model Stealing Attack Using Zeroth-Order Gradient Estimation [CVPR 2021] [S. Kariyappa](#), A. Prakash, M. K. Qureshi

Tolerating Noise in PCM-Based AI Accelerators via Noise-Aware Training [IEEE Transactions on Electron Devices 2021] [S. Kariyappa et al.](#)

Protecting DNNs from Theft using an Ensemble of Diverse Models [ICLR 2021] [S. Kariyappa](#), A. Prakash, M. K. Qureshi

Defending Against Model Stealing Attacks with Adaptive Misinformation [CVPR 2020] [S. Kariyappa](#), M. K. Qureshi

Improving Adversarial Robustness of Ensembles with Diversity Training [S. Kariyappa](#), M. K. Qureshi